

1 Inheritance

Consider the following example of simple inheritance in c++. It includes a derived class "Child" that inherits from the base class "Parent." It inherits Parent's x and y members and also has a public member of its own, "luckyFin."

```
1  class Parent
2  {
3      public:
4          void setX(int w)
5          {
6              x = w;
7          }
8          void setY(int h)
9          {
10             y = h;
11         }
12     protected:
13         int x;
14         int y;
15 };
16
17 class Child: public Parent
18 {
19     public:
20         int xPlusY()
21         {
22             return (x + y);
23         }
24
25         int luckyFin;
26 };
27
28 int main(void)
29 {
30     Child * nemo = new Child();
31
32     nemo->setX(2);
33     nemo->setY(3);
34
35     int add = nemo->xPlusY();
36
37     return 0;
38 }
```

The resulting assembly code is too long to list in full but snippets will be shown as the topics they are relevant to are discussed.

1.1 Initialization

The following x86 snippet shows the initialization of the nemo object.

```
1  call  _Znwj
2  add   esp, 16
3  mov   DWORD PTR [eax], 0
4  mov   DWORD PTR [eax+4], 0
5  mov   DWORD PTR [eax+8], 0
6  mov   DWORD PTR [ebp-12], eax
7  mov   eax, DWORD PTR [ebp-12]
```

The call to `_Znwj` returns a pointer to the location of the newly created nemo object in memory. Then the data members of nemo are moved into memory at that location.

1.2 Data Layout

From the code above it is not immediately clear which element is stored at which location. This information can be deduced by considering a call to `setX` inherited from the Parent class. The assembly for such a call is as follows.

```
1  push  2
2  push  eax
3  call  _ZN6Parent4setXEi
```

with the function itself being

```
1  _ZN6Parent4setXEi:
2  push  ebp
3  mov   ebp, esp
4  mov   eax, DWORD PTR [ebp+8]
5  mov   edx, DWORD PTR [ebp+12]
6  mov   DWORD PTR [eax], edx
7  pop   ebp
8  ret
```

Before the call is made, `eax` holds the address of the nemo object in memory. The `setX` method then moves the argument '2' located at `[ebp+12]` into `[eax]`, the lowest memory location in nemo. Similar analysis of the `setY` function shows us that `y` is located at `[eax+4]`. From this we can assume that 'luckyFin' is stored at `[eax+8]`.

The general rule that we can conclude from this examination is that within an object that inherits from a parent class the object's own member variables are stored first (at higher memory addresses) and then the inherited members are stored in reverse order with the first member of the parent class being stored at the lowest address.

1.3 Destruction

The first example will demonstrate what happens when a user created object goes out of scope. In order to more easily demonstrate this the c++ code has been modified in the following ways:

```
1 void scope() {
2     Child * nemo = new Child();
3     nemo->setX(2);
4     nemo->setY(3);
5     int add = nemo->xPlusY();
6 }
7
8 int main(void) {
9     scope();
10    return 0;
11 }
```

The last lines of assembly within the 'scope' function call look like this:

```
1     call _ZN5Child6xPlusYEv
2     add     esp, 16
3     mov     DWORD PTR [ebp-16], eax
4     leave
5     ret
```

As can be seen, the deallocation of the nemo object is very simple and operates just like deallocation of any other variables on the stack by adding to esp.

The second example will show what happens when something is 'destroyed' by the destructor.

```
1 class Planet {
2     public:
3         void setX(int w) {
4             x = w;
5         }
6         void setY(int h) {
7             y = h;
8         }
9     protected:
10        int x;
11        int y;
12 };
13
14 class DevourerOfWorlds {
15     public:
16         int luckyFin;
17         DevourerOfWorlds();
18         ~DevourerOfWorlds();
```

```
19
20     private:
21         Planet * earth;
22     };
23
24     DevourerOfWorlds::DevourerOfWorlds() {
25         earth = new Planet();
26     }
27
28     DevourerOfWorlds::~DevourerOfWorlds() {
29         delete earth;
30     }
31
32     void scope() {
33         DevourerOfWorlds galactus = DevourerOfWorlds();
34     }
35
36     int main(void) {
37         scope();
38         return 0;
39     }
```

Here the constructors and destructors have been explicitly defined. The destructor is automatically called when the program exits the scope of scope(). They compile to

```
1  _ZN16DevourerOfWorldsC2Ev:
2      push ebp
3      mov  ebp, esp
4      sub  esp, 8
5      sub  esp, 12
6      push 8
7      call _Znwj
8      add  esp, 16
9      mov  DWORD PTR [eax], 0
10     mov  DWORD PTR [eax+4], 0
11     mov  edx, DWORD PTR [ebp+8]
12     mov  DWORD PTR [edx+4], eax
13     leave
14     ret
15  _ZN16DevourerOfWorldsD2Ev:
16     push ebp
17     mov  ebp, esp
18     sub  esp, 8
19     mov  eax, DWORD PTR [ebp+8]
20     mov  eax, DWORD PTR [eax+4]
21     sub  esp, 12
22     push eax
23     call _ZdlPv
```

```
24     add    esp, 16
25     leave
26     ret
```

The constructor initializes galactus as we expect and also calls `_Znwj` to create the Planet earth object.

In the destructor, the address of earth is moved into `eax` which is then passed as an argument to `_ZdlPv` which deallocates the memory allocated to it. From there things proceed as normal with memory being deallocated by adding to `esp` and then `leave` and `ret` making up the standard epilogue.

2 Optimization

The following c++ code will be used to demonstrate a number of examples of different types of optimizations that might be done by a compiler. All optimizations are done using the `-O2` flag.

```
1  int constantArithmetic() {
2      int x = 0;
3      for (int i = 0; i < 100; i++) {
4          x += i;
5      }
6      return x;
7  }
8
9  int redundancy(int x) {
10     int y = 2 * x;
11     int z = 2 * x;
12     return y + z;
13 }
14
15 int simplification(int x) {
16     int y = x + x;
17     int z = 2 * x;
18     int ret = y - z;
19     return ret;
20 }
21
22 int inlining(int x) {
23     return x + 1;
24 }
25
26 int main(void) {
27     int arg;
28     cin >> arg;
29     int x = constantArithmetic();
30     int y = redundancy(arg);
31     int z = simplification(arg);
```

```
32     int a = inlining(arg);
33     cout << a;
34     return 0;
35 }
```

2.1 Constant Arithmetic

Compare the unoptimized version of the function

```
1 _Z18constantArithmeticv:
2     push ebp
3     mov  ebp, esp
4     sub  esp, 16
5     mov  DWORD PTR [ebp-4], 0
6     mov  DWORD PTR [ebp-8], 0
7     mov  eax, DWORD PTR [ebp-8]
8     add  DWORD PTR [ebp-4], eax
9     add  DWORD PTR [ebp-8], 1
10    cmp  DWORD PTR [ebp-8], 99
11    mov  eax, DWORD PTR [ebp-4]
12    leave
13    ret
```

to the optimized version.

```
1 _Z18constantArithmeticv:
2     mov  eax, 4950
3     ret
```

In the unoptimized version the code loops naively as we expect. The optimized version demonstrates the compiler's ability to do constant arithmetic at compile time. This function does not depend on any user input or variable quantities and will always return the same value. The compiler recognizes this and is able to compute the value, even though the calculation is done within a loop and isn't entirely trivial, and insert the value directly into the code.

2.2 Redundancy

Compare the unoptimized version of the function

```
1 _Z10redundancyi:
2     push ebp
3     mov  ebp, esp
4     sub  esp, 24
5     mov  eax, DWORD PTR [ebp+8]
6     add  eax, eax
7     mov  DWORD PTR [ebp-12], eax
8     mov  eax, DWORD PTR [ebp+8]
```

```
9      add    eax, eax
10     mov    DWORD PTR [ebp-16], eax
11     sub    esp, 8
12     push   DWORD PTR [ebp-12]
13     push   OFFSET FLAT:_ZSt4cout
14     call   _ZNSolsEi
15     add    esp, 16
16     sub    esp, 8
17     push   DWORD PTR [ebp-16]
18     push   OFFSET FLAT:_ZSt4cout
19     call   _ZNSolsEi
20     add    esp, 16
21     leave
22     ret
```

to the optimized version.

```
1 _Z10redundancyi:
2     push   ebx
3     sub    esp, 16
4     mov    eax, DWORD PTR [esp+24]
5     lea    ebx, [eax+eax]
6     push   ebx
7     push   OFFSET FLAT:_ZSt4cout
8     call   _ZNSolsEi
9     pop    eax
10    pop    edx
11    push   ebx
12    push   OFFSET FLAT:_ZSt4cout
13    call   _ZNSolsEi
14    add    esp, 24
15    pop    ebx
16    ret
```

The code in both cases is fairly long and involves strange function calls in order to print things to stdout. The optimization can be seen in lines 6-9 of the unoptimized code. It computes the values of z and z separately and then loads and prints each one individually. In the optimized version the compiler recognizes that y and z will have the same value. Thus it calculates the value once and then prints it twice, avoiding a redundant computation.

2.3 Simplification

Compare the unoptimized version of the function

```
1 _Z14simplificationi:
2     push   ebp
3     mov    ebp, esp
4     sub    esp, 16
```

```

5     mov     eax, DWORD PTR [ebp+8]
6     add     eax, eax
7     mov     DWORD PTR [ebp-4], eax
8     mov     eax, DWORD PTR [ebp+8]
9     add     eax, eax
10    mov     DWORD PTR [ebp-8], eax
11    mov     eax, DWORD PTR [ebp-4]
12    sub     eax, DWORD PTR [ebp-8]
13    mov     DWORD PTR [ebp-12], eax
14    mov     eax, DWORD PTR [ebp-12]
15    leave
16    ret

```

to the optimized version.

```

1  _Z14simplificationi:
2      xor     eax, eax
3      ret

```

The optimized version of the code recognizes that $x-x$ must equal 0 and simply places that value in `eax`. This is different from constant arithmetic because the value of x could be anything at runtime. The compiler is demonstrating awareness of basic mathematical axioms in order to simplify the result even when it knows nothing about the input.

2.4 Inlining

The difference here takes place not in the functions implementation but in its calling. In the unoptimized version of the code it is called as

```

1     push    eax
2     call    _Z8inliningi

```

In the optimized version no call is ever made to the function. The result of the function is displayed to the user so the compiler can't choose to just ignore it. Instead, recognizing that the function is trivially simple, it decided to inline the code into main. It looks like this.

```

1     add     eax, 1

```

In addition to the differences mentioned above there were some optimizations that took place in every function. First, the standard prologue and epilogue are ignored in the optimized version. Second, because the results of the simplification and constant arithmetic functions calls go unused neither of them were ever actually called in the optimized code.

References

I found the following links useful in trying to understand my assembly code.

1. http://en.wikipedia.org/wiki/Loop_unrolling
2. http://en.wikipedia.org/wiki/Redundant_Code
3. http://en.wikipedia.org/wiki/Category:Compiler_optimizations
4. <http://stackoverflow.com/questions/5474355/about-leave-in-x86-assembly>
5. http://www.compileroptimizations.com/category/expression_simplification.htm