

1 Contents of testfile4.txt

zzzzz zzzzy zzzzx zzzzw zzzzv zzzzu zzzzt zzzzs zzzzr zzzzq zzzzp
zzzzo zzzzn zzzzm zzzzl zzzzk zzzzj zzzzi zzzzh zzzzg zzzzf zzzze zzzzd
yyyyz yyyyy yyyyx yyyyw yyyyv yyyyu yyyyt yyyys yyyyr yyyyq
yyyyp yyyyoy yyyyn yyyym yyyyl yyyyk yyyyj yyyyi yyyyh yyyyg
yyyf yyye yyyd xxxz xxxxy xxxxx xxxxw xxxxv xxxxu xxxxt
xxxs xxxr xxxq xxxp xxxo xxxn xxxm xxxl xxxk xxxj
xxxi xxxh xxxg xxxf xxxe xxxd wwwwz wwwwy wwwwx
wwwvw wwwwv wwwwu wwwwt wwwws wwwwr wwwwq wwwwp
wwwwo wwwwn wwwwm wwwwl wwwwk wwwwj wwwwi wwwwh
wwwwg wwwwf wwwwe wwwwd vvvvz vvvvy vvvvx vvvvw vvvv
vvvvu vvvvt vvvvs vvvvr vvvvq vvvvp vvvvo vvvvn vvvvm vvvvl
vvvvk vvvvj vvvvi vvvvh vvvvg vvvvf vvvve vvvvd uuuuz uuuuy
uuuux uuuuw uuuuv uuuuu uuuut uuuus uuuur uuuuq uuuup uuuuo
uuuun uuuum uuuul uuuuk uuuuj uuuui uuuuh uuuug uuuuf uuuue
uuud ttttz tttty ttttx ttttw ttttv ttttu tttt tttts ttttr ttttq ttttp
tttto ttttn ttttm ttttl ttttk ttttj tttti tttth ttttg ttttf tttte ttttd ssssz
ssssy ssssx ssssw ssssv ssssu sssst sssss ssssr ssssq ssssp sssso ssssn
sssm sssl sssk sssj sssi sssh sssg sssf sssse ssssd rrrrz rrrry rrrrx
rrrrw rrrrv rrrru rrrrt rrrrs rrrr rrrrq rrrrp rrrro rrrrn rrrrm rrrrl
rrrk rrrj rrri rrrh rrrg rrrf rrrre rrrrd qqqqz qqqqy qqqqx qqqqw
qqqqv qqqqu qqqqt qqqqs qqqqr qqqq qqqqp qqqqo qqqqn qqqqm
qqqql qqqqk qqqqj qqqqi qqqqh qqqqg qqqqf qqqqe qqqqd ppppz
ppppy ppppx ppppw ppppv ppppu ppppt pppps ppppr ppppq ppppp
ppppo ppppn ppppm ppppl ppppk ppppj ppppi pppph ppppg ppppf
ppppe ppppd ooooz ooooy oooox oooow oooov oooou ooot oooos
oooor oooq ooop oooo ooon ooom oool oook oooj oooi
oooh ooog ooof oooe oood nnnnz nnnny nnnnx nnnnw nnnnv
nnnu nnnnt nnnns nnnnr nnnnq nnnnp nnnno nnnnn nnnnm nnnnl
nnnk nnnj nnni nnnh nnnng nnnnf nnnne nnnnd mmmmmz mm-
mmy mmmmx mmmmw mmmmv mmmmu mmmmt mmmms mm-
mmr mmmmq mmmmp mmmmo mmmmn mmmmm mmmml mm-
mmk mmmmj mmmmi mmmmh mmmmg mmmmf mmmme mm-
mmd llllz lllly llllx llllw llllv lllu lllt llls lllr lllq lllp llll
llllm lllk lllj llli lllh lllg lllf llll kkkkz kkkky kkkkx kkkkw
kkkkv kkkku kkkkt kkkks kkkkr kkkkq kkkkp kkkko kkkkn kkkkm
kkkl kkkk kkkkj kkkki kkkkh kkkkg kkkkf kkkke kkkkd jjjjz jjjjy
jjjx jjjw jjjv jjju jjjt jjjs jjjr jjjq jjjp jjjo jjjn jjjm jjjl jjjk jjjj
jjji jjjh jjjg jjjf jjje jjjd iiiz iiy iiix iiiv iiu iiit iiis iiir
iiiq iiip iiio iiii iiii iiii iiii iiii iiii iiii hhhh
hhhhx hhhh hhhh hhhh hhhh hhhh hhhh hhhh hhhh hhhh
hhhp hhhh hhhh hhhh hhhh hhhh hhhh hhhh hhhh hhhh
hhhhf hhhh hhhh hhhh ggggz ggggy ggggx ggggw ggggv ggggu ggggt
ggggs ggggr ggggq ggggp ggggo ggggn ggggm ggggl ggggk ggggj ggggi
ggggh ggggg ggggf gggge ggggd fffz fffy fffx fffv fffu ffft fffs

```

ffffr ffffq ffffp ffffo ffffn ffffm ffffl ffffk fffj fffi fffh fffg ffff fffe fffd
eeeez eeeey eeex eeew eeew eeew eeew eeew eeew eeew eeew eeew eeew eeew eeew
eeeo eeem eeem eeem eeem eeem eeem eeem eeem eeem eeem eeem eeem eeem eeem
eeed ddddz ddddy ddddx ddddw ddddv ddddu ddddt ddds dddr
dddq dddp dddo dddn dddm dddl dddk dddj dddi dddh
dddg dddf ddde dddd ant

```

2 Why it works

This testfile works because its contents are in reverse alphabetic order. This means that the element being inserted into the search tree is always lexicographically prior to previous elements and is inserted to the left of all existing nodes. In the unbalanced BST this results in a tree that is essentially equivalent to a singly linked list. Searching for the word “ant” demonstrates the difference in efficiency for worst case scenarios for a BST compared to an AVL Tree.

3 Numerical results

testfile1.txt - general statistics

Structure	Total Nodes	Average Node Depth	Single Rotations	Double Rotations
AVL	19	2.26316	6	2
BST	19	3.15789	n/a	n/a

testfile1.txt - search results

Word	AVL			BST		
	Left Links Followed	Right Links Followed	Total Links Followed	Left Links Followed	Right Links Followed	Total Links Followed
me	3	1	4	4	2	6
kind	1	0	1	3	1	4
problems	1	1	2	2	1	3
created	2	1	3	4	1	5

testfile2.txt - general statistics

Structure	Total Nodes	Average Node Depth	Single Rotations	Double Rotations
AVL	16	2.5	9	2
BST	16	6.0625	n/a	n/a

testfile2.txt - search results

Word	AVL			BST		
	Left Links Followed	Right Links Followed	Total Links Followed	Left Links Followed	Right Links Followed	Total Links Followed
caught	2	1	3	0	2	2
flying	1	1	2	0	5	5
flown	2	2	4	1	5	6
mauve	0	1	3	0	9	9

testfile3.txt - general statistics

Structure	Total Nodes	Average Node Depth	Single Rotations	Double Rotations
AVL	13	2.23077	5	0
BST	13	3.23077	n/a	n/a

testfile3.txt - search results

Word	AVL			BST		
	Left Links Followed	Right Links Followed	Total Links Followed	Left Links Followed	Right Links Followed	Total Links Followed
misty	2	1	3	3	2	5
cobwebs	1	0	1	1	0	0
clockwork	2	0	2	2	0	2
landscape	1	1	2	2	1	3

testfile4.txt - general statistics

Structure	Total Nodes	Average Node Depth	Single Rotations	Double Rotations
AVL	530	7.08868	520	2
BST	530	264.5	n/a	n/a

testfile4.txt - search results

Word	AVL			BST		
	Left Links Followed	Right Links Followed	Total Links Followed	Left Links Followed	Right Links Followed	Total Links Followed
ggggs	6	2	8	444	0	444
nnnni	3	4	7	293	0	293
uuuuo	6	2	8	126	0	126
ant	9	0	9	529	0	3

4 When are AVL trees preferable?

AVL trees are the better option when operations need to be reliably fast. This is not to imply that AVL trees are always faster than BSTs, but that they perform better in the worst case. In fact, due to the overhead created by rebalancing rotations in AVL trees, a BST will often be faster for a particular operation. Thus they are ideal when it is not critical that an operation is as fast as possible, but it is very important that any given operation will never take significantly longer than average. An example of when these types of goals might exist is when creating a user interface for an application. A page load or button press that takes much longer than normal would have a severe negative impact on the user experience.

Certain characteristics of the data being operated on and what is being used for can also influence the choice between BSTs and AVL trees. If the data is inserted into the tree in sorted order, a BST will simply mimic a linked list, provided none of the benefits of the tree data structure. If it is known that elements of the tree will be accessed with some uniformity then the reduced average path length of the AVL tree may be more important than if only a few key elements are being accessed repeatedly and can be stored near the root of the tree. Besides performance, the cost of implementation as discussed in the next section is also something to consider when deciding between the two tree types.

5 AVL implementation costs

Although they are generally faster, AVL trees are also much more difficult to implement. The addition of automatic balancing adds complexity to the code, including handling multiple different cases for inserting and removing nodes.

Additionally, to allow balancing to be done with relative speed, nodes must store additional information about their height in the tree. This means that AVL trees consume more memory than BSTs, an amount that grows in direct correspondance to the size of the tree.